


An AI agent deleted everything in 9 seconds

A Claude-powered coding agent. One API token. Zero guardrails.



02 ·

It wasn't the model that failed

The agent hit a credential mismatch in staging. Found a Railway API token with full GraphQL access. Used volumeDelete. Nine seconds. Production and backups — both gone.

80%

AI PROJECTS FAIL BEFORE PRODUCTION

3x

FASTER WITH PROPER INFRASTRUCTURE

0

TOLERANCE FOR FRAGILE SYSTEMS AT
VERYDIAN

Demo → Prod

WHERE MOST PROJECTS DIE

03 ·

Excessive permissions are always the first failure mode

The agent had broader access than any engineer would've been granted. No restrictions on destructive operations. No human confirmation required. That's not a model problem. It's an architecture problem.

80%

AI PROJECTS FAIL BEFORE PRODUCTION

3x

FASTER WITH PROPER INFRASTRUCTURE

0

TOLERANCE FOR FRAGILE SYSTEMS AT
VERYDIAN

Demo → Prod

WHERE MOST PROJECTS DIE

04 ·

The safety architecture gap is real — and growing

Companies are deploying agents into production faster than they're building systems to make those integrations safe. The PocketOS incident isn't an anomaly. It's what skipping the harness looks like.

80%

AI PROJECTS FAIL BEFORE PRODUCTION

3x

FASTER WITH PROPER INFRASTRUCTURE

0

TOLERANCE FOR FRAGILE SYSTEMS AT
VERYDIAN

Demo → Prod

WHERE MOST PROJECTS DIE

05 ·

Three rules before any agent touches production

Minimum permissions — exactly what the task needs.
Hard stops on irreversible actions. Full audit trails on every decision. These aren't constraints on capability. They're what production-ready looks like.

80%

AI PROJECTS FAIL BEFORE PRODUCTION

3x

FASTER WITH PROPER INFRASTRUCTURE

0

TOLERANCE FOR FRAGILE SYSTEMS AT
VERYDIAN


Demo → Prod

WHERE MOST PROJECTS DIE



Build the harness first.
Then the agent.

Eliminate operational gravity. Build systems that scale without scaling headcount.



VERYDIAN STACK