

↓ FRAMEWORK

Every AI agent failure has the same root cause

Not a bad model. Bad infrastructure around the model. The harness is the difference.

01 Your team copies data between systems daily

02 Decisions wait on a person, not data


03 You've automated tasks but not workflows

IF YES TO ANY →
You need an agent, not another app.

WHAT IT IS

2 The harness is what makes agents production-safe

It's the layer that transforms a capable model into a reliable, observable operation. Without it, you're deploying capability into production and hoping nothing unexpected happens. That's not a strategy.




SYS.NODES

THE COMPONENTS

3 Five components — in this order

Structured I/O. Timeout and resource limits. Full audit trails. Minimum permissions. Human confirmation gates on irreversible actions. Every production agent needs all five. Most demos have none.



COMPARISON

WHAT'S AVAILABLE NOW

4 Codex and Claude Code are building this in — slowly

Codex enterprise: sandbox policies, audit hooks, CA cert support. Claude Code: PostToolUse hooks for intercepting every tool call. The tooling is improving. But you can't wait for it to catch up.


DETECT	ROUTE	ACT
✓	—	—
—	✓	—
—	—	✓

DECISION.MAP

THE RISK

5 No harness means a production incident is a matter of when

Every agent you deploy without minimum permissions, audit logging, and human gates on destructive actions is a ticking clock. The PocketOS database deletion was nine seconds. Yours will take longer. Or it won't.



SYS.OUTPUT



SYS.FRAMEWORK / AGENTIC SYSTEMS

Build the harness first. Then build the agent.

Build agents that eliminate operational drag. Scale execution without scaling headcount.

