


171% ROI average. 86% see zero. Here's the gap.

The five root causes that separate production from pilot graveyard.


Root cause 1: Legacy data integration.

5–12 systems that don't talk cleanly. The agent hits a data wall in week one. No model solves a data pipeline problem.



Root cause 2: No monitoring for silent failures.

An agent producing subtly wrong outputs for two weeks causes compounding damage. Without monitoring from day one, you won't know until the damage is done.



DATA.PLOT


Root cause 3: Unclear ownership when it breaks.

Someone needs to own the fix in under 24 hours when an agent fails in production. If nobody knows whose problem it is, every hour compounds the cost.



Root causes 4 and 5: Volume and domain data.


Output quality that holds at 50 runs degrades at 50,000. Generic models produce generic outputs. Specialized industries need domain-grounded context, not generic capability.



The 14% that succeed scope infrastructure before the pilot.

Data pipelines ready. Monitoring live from day one.

Ownership defined before the first test run. That's not caution. That's why the math works.





EXEC.BRIEF / AI INFRASTRUCTURE



Follow me — I build in public

Build agents that belong in the 14%.