

```
$ verydian boot --topic "the failure data"
loading context ... done
```

88% of agent pilots
never reach production _

I checked the failure data. The model was never the problem.

RUN SESSION _____

```
arhem@verydian:~/the breakdown$ step 2
```

Three reasons, and none of them technical

Forrester traced the failures to unclear success criteria, insufficient tool access, and eval coverage that quietly decays. Not one model quality issue in the list.

STDOUT

01

CRITERIA 41%

02

ACCESS 33%

03

EVAL DRIFT 26%

```
arhem@verydian:~/cause one$ step 3
```

What clear success criteria actually requires

A single metric everyone agrees on. A named owner. A pass threshold defined before launch, not argued about after. A review date on the calendar.

STDOUT



DEFINED METRIC

PASS



NAMED OWNER

PASS



PASS THRESHOLD

LIVE



REVIEW CADENCE

PASS

```
arhem@verydian:~/cause two$ step 4
```

Where access requests quietly stall

The request gets approved. The scope gets granted for a sandbox. Then production data never actually arrives, and the pilot runs forever on fake inputs.

STDOUT

REQUEST
ACCESS

SCOPE GRANT

TEST SANDBOX

PRODUCTION
DATA

arhem@verydian:~/cause three\$ step 5

How eval coverage disappears without anyone noticing

Week one, every trace gets checked. Week four, it's spot checks. By week twelve, nobody's looking, and the drift has been compounding the whole time.




```
arhem@verydian:~/the whole system$ step 6
```

Everything has to be scoped around one build

Success criteria, data access, eval coverage, and ownership aren't separate checkboxes. Skip one and the other three don't hold either.

STDOUT





I define these three before an agent goes live

Every pilot I scope answers all three questions before it touches anything real.

FOLLOW FOR MORE