

CACHE ECONOMICS

Your bill depends on prompt order

Same model, same task. The difference is what sits at the top of the prompt.

UNROLL THE DRAWING

SECTION 02 · TWO PRICE LINES

Cached reads are priced separately

Claude Code 2.1.280 lists Opus 5.5 cache reads at \$0.20 per Mtok against \$4 for input. OpenAI offers a 90% discount on cached input reads.

FIG. 2

\$0.20 vs \$4

CACHE READ VS INPUT

● PER MTOK, OPUS 5.5

● OPENAI: 90% CACHED DISCOUNT

SCALE · PRODUCTION

VERYDIAN

SECTION 03 · PROMPT LAYOUT

Order the prompt by how often it changes

Stable instructions go first, tool definitions next, task data last. A cache hit needs an unchanged prefix.

FIG. 3

L3

TASK DATA (LAST)

L2

TOOL DEFINITIONS

L1

STABLE INSTRUCTIONS
(FIRST)

SECTION 04 · HIT VS MISS

One changed line breaks the prefix

The start of the prompt must match between calls. Change something early and the cached part no longer matches.

FIG. 4

CACHE HIT

- STABLE PREFIX
- VOLATILE DATA LAST

VS

CACHE MISS

- TIMESTAMP AT TOP
- NO CACHED READ

SECTION 05 · PREFIX BREAKERS

Where prefixes break unnoticed

Check these before blaming the model. Each one changes the start of the prompt between calls.

FIG. 5



TIMESTAMP IN SYSTEM PROMPT

LIVE



TOOL LIST REORDERED

PASS




PER-RUN ID AT THE TOP

PASS



VOLATILE DATA BEFORE RULES

PASS



Audit the top of your prompt.

Find anything that changes between calls and move it to the bottom.

REORDER THE PROMPT